

# Assignment Part-II

## Subjective Questions

### Question 1

Answer:

What is the optimal value of alpha for ridge and lasso regression? What will be the changes in the model if you choose double the value of alpha for both ridge and lasso? What will be the most important predictor variables after the change is implemented?

```
print(f"The optimal value for Ridge Regression is : {optimalvalue_ridge}")
print(f"The optimal value for Lasso Regression is : {optimalvalue_lasso}")
The optimal value for Ridge Regression is : 0.01
The optimal value for Lasso Regression is : 0.0001
# Doubling Lasso and Ridge Regression's alpha values
optimalvalue_ridge *= 2
optimalvalue_lasso *= 2
print(f"Doubled alpha values of Ridge is {optimalvalue_ridge} and Lasso is {optimalvalue_lasso}")
Doubled alpha values of Ridge is 0.02 and Lasso is 0.0002
```

### Build Lasso Regression model

```
alpha = optimalvalue_lasso
lasso = Lasso(alpha=alpha)
lasso.fit(X_train, y_train)
Lasso(alpha=0.0002)
lasso.coef_
array([-0.11816776,  0.40326807,  1.04310063,  0.20154269,  0.01320272,
        0.30796107,  0.3295398 , -0.28630072,  0.43099879,  0.1762002 ,
        0.46486274, -0.0752377 , -0.0096367 , -0.12095382,  0.1175943 ,
        0.1614835 ,  0.07220974,  0.12588766,  0.15207907,  0.11349623,
        0.16800439,  0.1053326 , -0.08579734, -0.01212515,  0.00742725,
       -0.23696381, -0.09273587, -0.1004028 ,  0.05524229, -0.05852677,
       -0.12689143, -0.22122537, -0.10950394,  0.          ,  0.14951854,
        0.06393483])
```

### Lasso features and their co-efficients

```
df_lasso = pd.DataFrame(index=X_train.columns)
df_lasso.rows = X_train.columns
df_lasso['Lasso'] = lasso.coef_
df_lasso
```

|            |           |
|------------|-----------|
|            | Lasso     |
| MSSubClass | -0.118168 |

|               |           |
|---------------|-----------|
|               | Lasso     |
| LotArea       | 0.403268  |
| OverallQual   | 1.043101  |
| OverallCond   | 0.201543  |
| BsmtUnfSF     | 0.013203  |
| BsmtFullBath  | 0.307961  |
| FullBath      | 0.329540  |
| KitchenAbvGr  | -0.286301 |
| TotRmsAbvGrd  | 0.430999  |
| Fireplaces    | 0.176200  |
| GarageCars    | 0.464863  |
| EnclosedPorch | -0.075238 |
| MSZoning_RH   | -0.009637 |
| LotShape_IR3  | -0.120954 |

|                      |           |
|----------------------|-----------|
|                      | Lasso     |
| LandContour_HLS      | 0.117594  |
| LandContour_Low      | 0.161483  |
| LandContour_Lvl      | 0.072210  |
| Neighborhood_Crawfor | 0.125888  |
| Neighborhood_NoRidge | 0.152079  |
| Neighborhood_NridgHt | 0.113496  |
| Neighborhood_StoneBr | 0.168004  |
| Neighborhood_Veenker | 0.105333  |
| BldgType_Twnhs       | -0.085797 |
| HouseStyle_2.5Unf    | -0.012125 |
| RoofStyle_Mansard    | 0.007427  |
| Exterior1st_BrkComm  | -0.236964 |
| Exterior1st_Stucco   | -0.092736 |

|                      |           |
|----------------------|-----------|
|                      | Lasso     |
| Exterior1st_Wd Sdng  | -0.100403 |
| Exterior2nd_Wd Sdng  | 0.055242  |
| BsmtQual_Fa          | -0.058527 |
| HeatingQC_Po         | -0.126891 |
| Functional_Maj2      | -0.221225 |
| Functional_Sev       | -0.109504 |
| SaleType_ConLD       | 0.000000  |
| SaleType_Oth         | 0.149519  |
| SaleCondition_Alloca | 0.063935  |

## Lasso Regression Model Evaluation

```

y_pred_train = lasso.predict(X_train)
y_pred_test = lasso.predict(X_test)

metric_double_1 = []
r2_train_lr = r2_score(y_train, y_pred_train)
print(f"Train r2 score is : {r2_train_lr}")
metric_double_1.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print(f"Test r2 score is : {r2_test_lr}")
metric_double_1.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print(f"Train RSS score is : {rss1_lr}")
metric_double_1.append(rss1_lr)

```

```

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print(f"Test RSS score is : {rss2_lr}")
metric_double_1.append(rss2_lr)

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print(f"Train MSE score is : {mse_train_lr}")
metric_double_1.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print(f"Test MSE score is : {mse_test_lr}")
metric_double_1.append(mse_test_lr**0.5)
Train r2 score is : 0.8514527457791814
Test r2 score is : 0.8328649333501988
Train RSS score is : 23.53222787609893
Test RSS score is : 12.418974514787804
Train MSE score is : 0.02302566328385414
Test MSE score is : 0.02835382309312284

```

### Build Ridge Regression model

```

alpha = optimalvalue_ridge
ridge = Ridge(alpha=alpha)
ridge.fit(X_train, y_train)

```

Out[180]:

```

Ridge(alpha=0.02)
ridge.coef_
array([-0.11453502,  0.49769529,  1.02880868,  0.20870043,  0.02323615,
        0.30571474,  0.32641994, -0.32622331,  0.44220508,  0.17049876,
        0.46661689, -0.07664049, -0.01803886, -0.15301168,  0.1355378 ,
        0.17556944,  0.08944529,  0.13511868,  0.15957298,  0.11897933,
        0.18028664,  0.13019827, -0.09222125, -0.02481312,  0.06377628,
       -0.33728585, -0.10829837, -0.11334924,  0.06975914, -0.06184982,
       -0.32566534, -0.27073409, -0.30979916,  0.02772084,  0.36717435,
        0.10859608])

```

### Ridge features and their co-efficients

```

df_ridge = pd.DataFrame(index=X_train.columns)
df_ridge.rows = X_train.columns
df_ridge['Ridge'] = ridge.coef_
df_ridge

```

|            | Ridge     |
|------------|-----------|
| MSSubClass | -0.114535 |
| LotArea    | 0.497695  |

|                 |           |
|-----------------|-----------|
|                 | Ridge     |
| OverallQual     | 1.028809  |
| OverallCond     | 0.208700  |
| BsmtUnfSF       | 0.023236  |
| BsmtFullBath    | 0.305715  |
| FullBath        | 0.326420  |
| KitchenAbvGr    | -0.326223 |
| TotRmsAbvGrd    | 0.442205  |
| Fireplaces      | 0.170499  |
| GarageCars      | 0.466617  |
| EnclosedPorch   | -0.076640 |
| MSZoning_RH     | -0.018039 |
| LotShape_IR3    | -0.153012 |
| LandContour_HLS | 0.135538  |

|                      |           |
|----------------------|-----------|
|                      | Ridge     |
| LandContour_Low      | 0.175569  |
| LandContour_Lvl      | 0.089445  |
| Neighborhood_Crawfor | 0.135119  |
| Neighborhood_NoRidge | 0.159573  |
| Neighborhood_NridgHt | 0.118979  |
| Neighborhood_StoneBr | 0.180287  |
| Neighborhood_Veenker | 0.130198  |
| BldgType_Twnhs       | -0.092221 |
| HouseStyle_2.5Unf    | -0.024813 |
| RoofStyle_Mansard    | 0.063776  |
| Exterior1st_BrkComm  | -0.337286 |
| Exterior1st_Stucco   | -0.108298 |
| Exterior1st_Wd Sdng  | -0.113349 |

|                      |           |
|----------------------|-----------|
|                      | Ridge     |
| Exterior2nd_Wd Sdng  | 0.069759  |
| BsmtQual_Fa          | -0.061850 |
| HeatingQC_Po         | -0.325665 |
| Functional_Maj2      | -0.270734 |
| Functional_Sev       | -0.309799 |
| SaleType_ConLD       | 0.027721  |
| SaleType_Oth         | 0.367174  |
| SaleCondition_Alloca | 0.108596  |

## Ridge Regression Model Evaluation

```

y_pred_train = ridge.predict(X_train)
y_pred_test = ridge.predict(X_test)

metric_double_r = []
r2_train_lr = r2_score(y_train, y_pred_train)
print(f"Train r2 score is : {r2_train_lr}")
metric_double_r.append(r2_train_lr)

r2_test_lr = r2_score(y_test, y_pred_test)
print(f"Test r2 score is : {r2_test_lr}")
metric_double_r.append(r2_test_lr)

rss1_lr = np.sum(np.square(y_train - y_pred_train))
print(f"Train RSS score is : {rss1_lr}")
metric_double_r.append(rss1_lr)

rss2_lr = np.sum(np.square(y_test - y_pred_test))
print(f"Test RSS score is : {rss2_lr}")
metric_double_r.append(rss2_lr)

```



```

mse_train_lr = mean_squared_error(y_train, y_pred_train)
print(f"Train MSE score is : {mse_train_lr}")
metric_double_r.append(mse_train_lr**0.5)

mse_test_lr = mean_squared_error(y_test, y_pred_test)
print(f"Test MSE score is : {mse_test_lr}")
metric_double_r.append(mse_test_lr**0.5)
Train r2 score is : 0.8530992328545421
Test r2 score is : 0.8277656774325617
Train RSS score is : 23.271398355851826
Test RSS score is : 12.797874829095916
Train MSE score is : 0.022770448489091807
Test MSE score is : 0.029218892303871955

```

## Comparison of co-efficients after Regularization

```

comparison['Ridge_Double'] = ridge.coef_
comparison['Lasso_Double'] = lasso.coef_

comparison.sort_values(by='Lasso', ascending=False)

```

|              | Linear   | Ridge    | Lasso    | Ridge_Double | Lasso_Double |
|--------------|----------|----------|----------|--------------|--------------|
| OverallQual  | 1.030270 | 1.029538 | 1.036706 | 1.028809     | 1.043101     |
| GarageCars   | 0.466611 | 0.466613 | 0.465674 | 0.466617     | 0.464863     |
| LotArea      | 0.502954 | 0.500309 | 0.453153 | 0.497695     | 0.403268     |
| TotRmsAbvGrd | 0.442432 | 0.442320 | 0.436777 | 0.442205     | 0.430999     |
| FullBath     | 0.325870 | 0.326146 | 0.327681 | 0.326420     | 0.329540     |
| BsmtFullBath | 0.305247 | 0.305484 | 0.306589 | 0.305715     | 0.307961     |
| SaleType_Oth | 0.375355 | 0.371221 | 0.262360 | 0.367174     | 0.149519     |

|                      | Linear   | Ridge    | Lasso    | Ridge_Double | Lasso_Double |
|----------------------|----------|----------|----------|--------------|--------------|
| OverallCond          | 0.208632 | 0.208667 | 0.204988 | 0.208700     | 0.201543     |
| Neighborhood_StoneBr | 0.180237 | 0.180262 | 0.174085 | 0.180287     | 0.168004     |
| Fireplaces           | 0.169985 | 0.170243 | 0.173061 | 0.170499     | 0.176200     |
| LandContour_Low      | 0.175476 | 0.175524 | 0.168432 | 0.175569     | 0.161483     |
| Neighborhood_NoRidge | 0.159484 | 0.159529 | 0.155763 | 0.159573     | 0.152079     |
| Neighborhood_Crawfor | 0.135249 | 0.135184 | 0.130558 | 0.135119     | 0.125888     |
| LandContour_HLS      | 0.135634 | 0.135586 | 0.126574 | 0.135538     | 0.117594     |
| Neighborhood_Veenker | 0.130385 | 0.130292 | 0.117864 | 0.130198     | 0.105333     |
| Neighborhood_NridgHt | 0.118791 | 0.118885 | 0.116130 | 0.118979     | 0.113496     |
| SaleCondition_Alloca | 0.109234 | 0.108913 | 0.086547 | 0.108596     | 0.063935     |
| LandContour_Lvl      | 0.089749 | 0.089597 | 0.080940 | 0.089445     | 0.072210     |
| Exterior2nd_Wd Sdng  | 0.070046 | 0.069902 | 0.062630 | 0.069759     | 0.055242     |
| RoofStyle_Mansard    | 0.064157 | 0.063966 | 0.035776 | 0.063776     | 0.007427     |

|                     | Linear    | Ridge     | Lasso     | Ridge_Double | Lasso_Double |
|---------------------|-----------|-----------|-----------|--------------|--------------|
| BsmtUnfSF           | 0.022946  | 0.023092  | 0.018083  | 0.023236     | 0.013203     |
| SaleType_ConLD      | 0.027801  | 0.027761  | 0.011742  | 0.027721     | 0.000000     |
| MSZoning_RH         | -0.017834 | -0.017937 | -0.013784 | -0.018039    | -0.009637    |
| HouseStyle_2.5Unf   | -0.024813 | -0.024813 | -0.018555 | -0.024813    | -0.012125    |
| BsmtQual_Fa         | -0.061655 | -0.061753 | -0.060121 | -0.061850    | -0.058527    |
| EnclosedPorch       | -0.076377 | -0.076510 | -0.075749 | -0.076640    | -0.075238    |
| BldgType_Twnhs      | -0.092248 | -0.092235 | -0.089108 | -0.092221    | -0.085797    |
| Exterior1st_Stucco  | -0.108609 | -0.108453 | -0.100692 | -0.108298    | -0.092736    |
| Exterior1st_Wd Sdng | -0.113638 | -0.113493 | -0.107008 | -0.113349    | -0.100403    |
| MSSubClass          | -0.114461 | -0.114498 | -0.116211 | -0.114535    | -0.118168    |
| LotShape_IR3        | -0.154117 | -0.153562 | -0.137553 | -0.153012    | -0.120954    |
| Functional_Sev      | -0.316220 | -0.312977 | -0.212884 | -0.309799    | -0.109504    |
| HeatingQC_Po        | -0.332267 | -0.328933 | -0.229611 | -0.325665    | -0.126891    |

|                     | Linear    | Ridge     | Lasso     | Ridge_Double | Lasso_Double |
|---------------------|-----------|-----------|-----------|--------------|--------------|
| Functional_Maj2     | -0.272084 | -0.271407 | -0.246685 | -0.270734    | -0.221225    |
| Exterior1st_BrkComm | -0.340590 | -0.338929 | -0.288809 | -0.337286    | -0.236964    |
| KitchenAbvGr        | -0.327021 | -0.326620 | -0.306698 | -0.326223    | -0.286301    |

## Comparison of metrics after Regularization

```
rg_metric = pd.Series(metric_double_r, name = 'Double Ridge Regression')
ls_metric = pd.Series(metric_double_l, name = 'Double Lasso Regression')

final_metric = pd.concat([final_metric, rg_metric, ls_metric], axis = 1)

final_metric
```

|   | Metric           | Linear Regression | Ridge Regression | Lasso Regression | Double Ridge Regression | Double Lasso Regression |
|---|------------------|-------------------|------------------|------------------|-------------------------|-------------------------|
| 0 | R2 Score (Train) | 0.853101          | 0.853100         | 0.852686         | 0.853099                | 0.851453                |
| 1 | R2 Score (Test)  | 0.827597          | 0.827683         | 0.830786         | 0.827766                | 0.832865                |
| 2 | RSS (Train)      | 23.271138         | 23.271204        | 23.336857        | 23.271398               | 23.532228               |
| 3 | RSS (Test)       | 12.810372         | 12.804017        | 12.573482        | 12.797875               | 12.418975               |
| 4 | MSE (Train)      | 0.150898          | 0.150898         | 0.151111         | 0.150899                | 0.151742                |

|   | Metric     | Linear Regression | Ridge Regression | Lasso Regression | Double Ridge Regression | Double Lasso Regression |
|---|------------|-------------------|------------------|------------------|-------------------------|-------------------------|
| 5 | MSE (Test) | 0.171019          | 0.170976         | 0.169430         | 0.170935                | 0.168386                |

## Question 2

You have determined the optimal value of lambda for ridge and lasso regression during the assignment. Now, which one will you choose to apply and why?

Answer:

final\_metric

|   | Metric           | Linear Regression | Ridge Regression | Lasso Regression | Double Ridge Regression | Double Lasso Regression |
|---|------------------|-------------------|------------------|------------------|-------------------------|-------------------------|
| 0 | R2 Score (Train) | 0.853101          | 0.853100         | 0.852686         | 0.853099                | 0.851453                |
| 1 | R2 Score (Test)  | 0.827597          | 0.827683         | 0.830786         | 0.827766                | 0.832865                |
| 2 | RSS (Train)      | 23.271138         | 23.271204        | 23.336857        | 23.271398               | 23.532228               |
| 3 | RSS (Test)       | 12.810372         | 12.804017        | 12.573482        | 12.797875               | 12.418975               |
| 4 | MSE (Train)      | 0.150898          | 0.150898         | 0.151111         | 0.150899                | 0.151742                |
| 5 | MSE (Test)       | 0.171019          | 0.170976         | 0.169430         | 0.170935                | 0.168386                |

### Question 3

After building the model, you realised that the five most important predictor variables in the lasso model are not available in the incoming data.

You will now have to create another model excluding the five most important predictor variables. Which are the five most important predictor variables now?

Answer:

#### Looking at the top 5 important predictor variables in Lasso model

```
comparison.sort_values(by='Lasso',ascending=False).head()
```

|              | Linear   | Ridge    | Lasso    | Ridge_Double | Lasso_Double |
|--------------|----------|----------|----------|--------------|--------------|
| OverallQual  | 1.030270 | 1.029538 | 1.036706 | 1.028809     | 1.043101     |
| GarageCars   | 0.466611 | 0.466613 | 0.465674 | 0.466617     | 0.464863     |
| LotArea      | 0.502954 | 0.500309 | 0.453153 | 0.497695     | 0.403268     |
| TotRmsAbvGrd | 0.442432 | 0.442320 | 0.436777 | 0.442205     | 0.430999     |
| FullBath     | 0.325870 | 0.326146 | 0.327681 | 0.326420     | 0.329540     |

```
# Looking at the current top 5 important predictor variables in Lasso model
comparison.sort_values(by='Lasso',ascending=False).Lasso.head(5)
```

```
OverallQual    1.036706
GarageCars      0.465674
LotArea         0.453153
TotRmsAbvGrd   0.436777
FullBath        0.327681
Name: Lasso, dtype: float64
```

```
# Creating a list to hold the current top 5 important predictor variables
top5_names = list(comparison['Lasso'].sort_values(ascending=False).head(5).index)
top5_names
['OverallQual', 'GarageCars', 'LotArea', 'TotRmsAbvGrd', 'FullBath']
# Drop the top 5 important predictor variables from X_train and X_test
X_train = X_train.drop(top5_names, axis=1)
```

```

X_test = X_test.drop(top5_names, axis=1)

print(X_train.shape)
print(X_test.shape)
(1022, 31)
(438, 31)

# List of alphas to tune

params = { 'alpha' : [0.0001,0.001,0.01,0.05,0.1,0.2,0.3,0.4,0.5,0.6,
                      0.7,0.8,0.9,1.0,2.0,3.0,4.0,5.0,6.0,7.0,
                      8.0,9.0,10.0,20,50,100,500,1000]}

# Applying lasso regression with 5 fold cross validation

lasso = Lasso()
folds = 5
model_cv = GridSearchCV(estimator=lasso,
                        param_grid=params,
                        scoring='neg_mean_absolute_error',
                        cv=folds,
                        return_train_score=True,
                        verbose=1)
model_cv.fit(X_train, y_train)
Fitting 5 folds for each of 28 candidates, totalling 140 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent worker
s.
[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed: 2.8s finished

GridSearchCV(cv=5, estimator=Lasso(),
             param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
                                   0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                                   4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
                                   100, 500, 1000]},
             return_train_score=True, scoring='neg_mean_absolute_error',
             verbose=1)

cv_results.shape
Out[196]:
(28, 21)

# Plotting train scores with alpha

plt.plot(cv_results['param_alpha'], cv_results['mean_train_score'], color='green')
plt.xlabel('alpha')
plt.ylabel("Negative Mean Absolute Error")

plt.title("Neg MAE and Alphas")
plt.show()

# Plotting testing scores with alpha
plt.plot(cv_results['param_alpha'], cv_results['mean_test_score'], color='orange')
plt.xlabel('alpha')
plt.ylabel("Negative Mean Absolute Error")

```

```
plt.title("Neg MAE and Alphas")
plt.show()
```

## Getting the optimal value of lambda

```
optimalvalue_lasso = model_cv.best_params_['alpha']
optimalvalue_lasso
Out[199]:
```

```
0.0001
```

## Build final Lasso Regression model

```
alpha = optimalvalue_lasso
lasso = Lasso(alpha=alpha)
lasso.fit(X_train, y_train)
```

```
Lasso(alpha=0.0001)
```

```
lasso.coef_
```

```
array([ 0.00228341,  0.21978345,  0.55040169,  0.50982786, -0.32405959,
        0.5890991 , -0.21277777, -0.05180727,  0.08867101,  0.29104398,
        0.18345996,  0.12195927,  0.10987309,  0.48620204,  0.43943659,
        0.42136982,  0.19854878, -0.18712857, -0.          ,  0.0650244 ,
       -0.56986175, -0.07587302, -0.12890651, -0.00653112, -0.15232127,
       -0.28407441, -0.52231084, -0.03715972, -0.          ,  0.00847051,
        0.22868149])
```

## Lasso features and their co-efficients

```
df_lasso = pd.DataFrame(index=X_train.columns)
df_lasso.rows = X_train.columns
df_lasso['Lasso'] = lasso.coef_
```

## Getting the new top 5 important predictor variables via Lasso Regression

```
df_lasso.sort_values(by='Lasso', ascending=False).head(5)
```

|            |          |
|------------|----------|
|            | Lasso    |
| Fireplaces | 0.589099 |



|                      |          |
|----------------------|----------|
|                      | Lasso    |
| BsmtUnfSF            | 0.550402 |
| BsmtFullBath         | 0.509828 |
| Neighborhood_NoRidge | 0.486202 |
| Neighborhood_NridgHt | 0.439437 |

## Question 4

How can you make sure that a model is robust and generalisable? What are the implications of the same for the accuracy of the model and why?

### Answer:

A robust model has low variance. This means that an unprecedented change in one or more features does not significantly alter the value of the predicted variable. Similarly, a generalizable model has reduced model complexity. As the number of features increase in the model, it becomes more complex which usually leads to low bias but high variance. A generalizable model has just enough features that it has as much low variance as possible.

This can be observed from the Bias-Variance tradeoff visual shown below.

The OLS (Ordinary least squares) regression model is very sensitive to outliers and they induce high variance. To reduce this, we can go ahead with regularization (Ridge/Lasso) which include a penalty term in the cost function of the model. This penalty term will move the coefficients of the model towards 0 and thus it reduces model complexity (as feature addition is heavily discouraged). This reduces overfitting in the model.

So regularization gets us high variance with a small trade-off in bias. Thus it helps us build a model which is robust and generalizable. A robust and generalizable model will have a good, consistent train as well as test accuracy.